

Movie review classification

- Major NLP task: sentiment classification
 - Work out the sentiment behind a piece
 - Some problems about target:
 - God damnit I dropped my iPhone
 - How does the author feel?
 - ... about their iPhone
- We're going to be crude: +ve and -ve only
 - You must have an opinion!



Experimental setup

- Read in the corpus
- This is going to contain many words! $|V| \rightarrow$ a lot
- Let's get just the k most frequent words
 - 2000 is a nice number to start with
- Based on these, write a feature extractor that includes only popular words
 - Use `set()` to remove dupes

Test feature extraction

```
>>> print( document_features( movie_reviews.words( 'pos/cv957_8737.txt' )))
```

```
{'contains(waste)': False, 'contains(lot)': False, ...}
```

- So far, so good
- Let's evaluate our system
 - Build a test/train split
 - Train a classifier
 - Evaluate on the test split
- Use the `.accuracy()` method

What words indicate review type?

- We can look at the most informative features
- This gives words that indicate +ve/-ve reviews
- Most Informative Features

contains(outstanding) = True	pos : neg	=	11.1 : 1.0
contains(seagal) = True	neg : pos	=	7.7 : 1.0
contains(wonderfully) = True	pos : neg	=	6.8 : 1.0
contains(damon) = True	pos : neg	=	5.9 : 1.0
contains(wasted) = True	neg : pos	=	5.8 : 1.0

- Oh dear, Steven...
- What happens if we select a lot more words as features?

Data mining: N-grams

- A secondary use of this classifier:
 - Finding indicative patterns in the data
- “Seagal” could be any Seagal (probably isn't..)
- What if we knew which one it was?
 - New feature extraction code
 - Find all n-grams, for 1...k n-grams

Data mining: n-grams

```
from nltk.util import ngrams  
.. everygrams
```

- See: <http://www.nltk.org/api/nltk.html>
- The goal is to convert:
 - [you, ai, n't, no, muslim, bruv]
 - [you, ai], [ai, n't], [n't, no], [no, muslim], [muslim, bruv]

