

# Introduction

- Optimisation technique
- Popular for tuning neural networks
- Previously used in MaxEnt, CRF, ...
- Goal:
  - Minimise an objective function  $J(\theta)$
  - $\theta \in \mathbb{R}^d$  are a model's parameters
  - Update params in opposite direction to gradient  $\nabla_{\theta} J$
  - Learning rate  $\eta$  determines step size

# Terminology

- $x$  : examples
- $y$  : labels
- $J$  : loss function (how bad the mistakes are)  
e.g. training error; the error on the training data
- $\theta$  : parameters
- $\eta$  : rate

# Three variants

- Batch gradient descent
  - Vanilla, original
  - Compute loss over whole training dataset at a time
- Stochastic gradient descent
  - Loss measured per example & label
- Mini-batch gradient descent
  - “Best of both worlds”
  - Loss computed over subset of training dataset

# Batch gradient descent

- Computes gradient of cost function w.r.t. to parameters  $\theta$  for the entire training dataset:
  - $\theta = \theta - \eta \nabla_{\theta} J(\theta)$
- Need to calculate the gradients for the whole dataset to perform just one update
  - Batch gradient descent can be slow and is intractable for datasets that don't fit in memory
  - Batch gradient descent also doesn't allow us to update our model online, i.e. with new examples on-the-fly.

# Batch gradient descent

- Code might look like
  - `for i in range(nb_epochs):`
  - `params_grad = evaluate_gradient(loss_function, data, params)`
  - `params = params - learning_rate * params_grad`
- Given a set number of training epochs:
  - Compute the gradient vector of the loss function for the whole dataset, w.r.t. our parameter vector `params`
  - Update our parameters in the opposite direction of the gradients
  - Learning rate  $\eta$  determines update size

# Batch gradient descent


- Guaranteed to:
  - Converge to the **global minimum for convex** error surfaces
  - Converge to a **local minimum for non-convex** surfaces.

# Stochastic gradient descent

- Stochastic = Random
- SGD performs a parameter update for each training example  $x^{(i)}$  and label  $y^{(i)}$ :
  - $\theta = \theta - \eta \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$
- Batch GD does redundant work:
  - Recomputes gradients for similar examples before doing parameter update
- SGD more efficient:
  - One update at a time
  - What advantage does this give?
  - Possible to apply to online learning

# Stochastic gradient descent

- SGD performs frequent updates:
  - high variance
  - Causes objective function to fluctuate heavily





# Stochastic gradient descent

- Other advantages of SGD vs. batch?
  - enables it to jump to new and potentially better local minima
- Side-effect of SGD's approach to finding minima?
  - Complicates convergence: SGD will keep overshooting
- Solution for this?
  - Slowly decrease learning rate

# Stochastic gradient descent

- Example code:
  - `for i in range(nb_epochs):`
  - `np.random.shuffle(data)`
  - `for example in data:`
  - `params_grad = evaluate_gradient(loss_function, example, params)`
  - `params = params - learning_rate * params_grad`
- Similar behaviour to batch GD:
  - Almost certainly converging to a **local minimum for non-convex** optimisation
  - Almost certainly converging to a **global minimum for convex** optimisation

# Mini-batch gradient descent

- Aims to take best of both worlds
- Performs update for every mini-batch of  $n$  examples:
  - $\theta = \theta - \eta \nabla_{\theta} J(\theta; x^{(i:i+n)}, y^{(i:i+n)})$
- Reduces variance of parameter updates
  - Leads to more stable convergence
- Can use some neat optimisations
  - E.g. GPU matrix factorisation
- Algorithm of choice for NN training

# Mini-batch gradient descent

- For  $n=50$ :
  - for  $i$  in range(nb\_epochs):
  - `np.random.shuffle(data)`
  - for batch in get\_batches(data, batch\_size=50):
  - `params_grad = evaluate_gradient(loss_function, batch, params)`
  - `params = params - learning_rate * params_grad`
- This updates parameters based on a  $J$  measured from a random subset of 50 examples

# Mini-batch gradient descent

- Does not guarantee good convergence
- Choosing a proper learning rate can be difficult
  - Too small: very slow convergence
  - Too high:
    - Damages convergence
    - Loss rate can fluctuate or even diverge
- Learning rate “schedules” reduce  $\eta$  at preset times, or when objective delta falls below threshold

# Mini-batch gradient descent


- Not all features have the same variance
  - A universal learning rate doesn't make sense here
  - Do bigger updates for sparser features?
- Saddle points are a challenge
  - Neural nets tend to have many local minima
  - These tend to meet in saddle points
  - Equates to a plateau of the same error
  - Hard to escape using gradient-based method, as gradient  $\rightarrow 0$  in all dimensions

# Gradient descent optimization

- Momentum
- Adagrad
- Adadelata
- RMSprop

# Momentum optimisation

- SGD has trouble navigating “ravines”
  - areas where the surface curves much more steeply in one dimension than in another
  - Ravines more common than local optima
- SGD will oscillate across slopes while making poor progress to the optimum





# Momentum optimisation


- Helps accelerate SGD in relevant direction
- Dampens oscillations
- Achieved by adding part  $\gamma$  of the last update to the current update
  - $v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$
  - $\theta = \theta - v_t$
- Parameter  $\gamma$  is the momentum

# Momentum optimisation

- Using momentum is like pushing a ball down a hill.
  - The ball accelerates as it rolls downhill, becoming faster and faster..
  - ..until it reaches its terminal velocity if there is air resistance, i.e.  $\gamma < 1$
- The momentum term:
  - increases for dimensions whose gradients point in the same directions
  - reduces updates for dimensions whose gradients change directions.
- This gives faster convergence & reduced oscillation.

# Momentum optimisation

- An example of momentum optimisation
  - Slows changes in high-gradient dimensions
  - Accelerates change in shallower dimensions



- Now we adapt to error function, why not adapt to parameter importance?

# Adam optimisation

- Adaptive Moment Estimation (Adam) computes adaptive learning rates for each parameter.
- It stores two factors:
  - exponentially decaying average of past squared gradients  $v_t$
  - exponentially decaying average of past gradients  $m_t$ , similar to momentum

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

# Adam optimisation

- $m_t$  and  $v_t$  are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients
- $m_t$  and  $v_t$  are initialized as vectors of 0's
  - they are biased towards zero,
  - especially during the initial time steps,
  - especially when the decay rates are small (i.e.  $\beta_1$  and  $\beta_2$  are close to 1).
- How can we avoid this?
  - Correct for bias in  $m$  and  $v$

# Adam optimisation

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}.$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

- Big beta will magnify  $m$  and  $v$

# Adam optimisation

- The adam update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

- Typical values:
  - 0.9 for  $\beta_1$ , 0.999 for  $\beta_2$ , and  $10^{-8}$  for  $\epsilon$

# Visualisation of algorithms

- <http://i.imgur.com/VkTJV VX.gif>
- Adagrad, Adadelata, and RMSprop almost immediately head off in the right direction and converge similarly fast
- Momentum and NAG are led off-track, evoking the image of a ball rolling down the hill.



# Visualisation at a saddle

- <http://i.imgur.com/1Awcohc.gif>
- SGD, Momentum, and NAG find it difficult to break symmetry, although the two latter eventually manage to escape the saddle point
- Adagrad, RMSprop, and Adadelta quickly head down the negative slope